

Performance

- What does it mean when you say “system A is of a higher performance as compared to system B”?

Performance

- What does it mean when you say “system A is of a higher performance as compared to system B”?
- “Time taken to perform a certain task T on system A is less than that taken to perform the same task T on system B.”
 - Note that this statement is specific to the task T. For some other task T', it may be possible that system B is the better performer.

Performance

```
int f = 1;
for(int i=2; i<=n; i++)
{
    f = f * i;
}
cout << f;
```

Program A

```
int f = 1;
for(int i=2; i<=2*n; i++)
{
    if (i <= n)
        f = f * i;
}
cout << f;
```

Program B

- System A: Program A is run on processor type P
- System B: Program B is run on processor type P
- Which system displays higher performance?

Performance

```
.main:
[0]      mov r2, 3
[4]      mov r1, 1
.loop:
[8]      mul r1, r1, r2
[12]     sub r2, r2, 1
[16]     cmp r2, 1
[20]     bgt .loop
[24]     st r1, 4[r0]
[28]     end
```

How many instructions are executed?

Performance

```
.main:
[0]      mov r2, 3
[4]      mov r1, 1
.loop:
[8]      mul r1, r1, r2
[12]     sub r2, r2, 1
[16]     cmp r2, 1
[20]     bgt .loop
[24]     st r1, 4[r0]
[28]     end
```

- Number of static instructions = 8
- Number of dynamic instructions = 12

Performance

```
int f = 1;
for(int i=2; i<=n; i++)
{
    f = f * i;
}
cout << f;
```

Program A

```
int f = 1;
for(int i=2; i<=2*n; i++)
{
    if (i <= n)
        f = f * i;
}
cout << f;
```

Program B

- System A: Program A is run on processor type P
- System B: Program B is run on processor type P
- System A displays higher performance because program A has fewer dynamic instructions (assuming every instruction takes equal amount of time)

SimpleRISC Processor Frequency

- How long does it take to execute a single instruction on our SimpleRISC processor?

SimpleRISC Processor Frequency

- How long does it take to execute a single instruction on our SimpleRISC processor?
- Work out the longest path in the combinational circuitry. Processor clock period should be greater than the time taken to execute this longest path. Let this clock period be t seconds. Processor frequency $f = 1/t$.
- Every instruction is executed in t s.
- Every instruction is executed in 1 cycle in our SimpleRISC processor.
 - Cycles per instruction, $CPI = 1$
 - Instructions per cycle, $IPC = 1$

Computing the Time a Program Takes

$$\begin{aligned}\tau &= \text{\#seconds} \\ &= \frac{\text{\#seconds}}{\text{\#cycles}} * \frac{\text{\#cycles}}{\text{\#instructions}} * (\text{\#instructions}) \\ &= \underbrace{\frac{\text{\#seconds}}{\text{\#cycles}}}_{1/f} + \underbrace{\frac{\text{\#cycles}}{\text{\#instructions}}}_{CPI} * (\text{\#instructions}) \\ &= \frac{CPI * \text{\#insts}}{f}\end{aligned}$$

- * CPI \rightarrow Cycles per instruction
- * f \rightarrow frequency (cycles per second)

The Performance Equation

$$P \propto \frac{IPC * f}{\#insts}$$

- * **IPC** → 1/CPI (Instructions per Cycle)
- * What are the **units** of performance ?
 - * **ANSWER** : arbitrary

Number of Instructions (#insts)

Static Instruction: The **binary** or executable of a **program**, contains a list of *static instructions*.

Dynamic Instruction: A **dynamic instruction** is a running instance of a static instruction, which is created by the **processor** when an instruction enters the pipeline.

- * Note that these are **dynamic** instructions
 - * **NOT** static instructions
- * A smart compiler can reduce the number of executed instructions

Number of Instructions(#insts) – 2

- * Dead code removal

- * Often **programmers** write **code** that does not determine the final output
- * This code is **redundant**
- * It can be identified and removed by the **compiler**


- * Function inlining

- * Very small **functions** have a lot of **overhead** → call, ret instructions, register spilling, and restoring
- * Paste the code of the **callee** in the code of the **caller** (known as **inlining**)

Computer Organisation and Architecture

**Smruti Ranjan Sarangi,
IIT Delhi**


Chapter 9 Principles of Pipelining



These slides are meant to be used along with the book: Computer Organisation and Architecture, Smruti Ranjan Sarangi, McGrawHill 2015
Visit: <http://www.cse.iitd.ernet.in/~srsarangi/archbooksoft.html>

Outline

- * Overview of Pipelining
- * A Pipelined Data Path
- * Pipeline Hazards
- * Pipeline with Interlocks
- * Forwarding
- * Performance Metrics
- * Interrupts/ Exceptions



Up till now

- * We have designed a **processor** that can execute all the **SimpleRisc** Instructions
- * We have look at two **styles** :
 - * With a **hardwired** control unit
 - * **Microprogrammed** control unit
 - * Microprogrammed data path
 - * Microassembly Language
 - * Microinstructions


Designing Efficient Processors

- * **Microprogrammed** processors are much slower than **hardwired** processors
- * Even **hardwired** processors
 - * Have a lot of **waste** !!!
 - * We have 5 stages.
 - * What is the IF stage doing, when the MA stage is active ?
 - * **ANSWER** : It is **idling**

The Notion of Pipelining

- * Let us go back to the **car assembly line**
 - * Is the **engine shop** idle, when the **paint shop** is painting a car ?
 - * **NO** : It is building the engine of another car
 - * When this engine goes to the body shop, it builds the engine of another car, and **so on**
- * **Insight** :
 - * **Multiple cars** are built at the same time.
 - * A car **proceeds** from one stage to the next

Pipelined Processors



- * The IF, ID, EX, MA, and RW stages process 5 instructions simultaneously
- * Each instruction proceeds from one stage to the next
- * This is known as pipelining

Advantages of Pipelining

- * We keep all parts of the data path, busy all the time
- * Let us assume that all the 5 stages do the same amount of **work**
 - * **Without** pipelining, every **T** seconds, an instruction completes its **execution**
 - * **With** pipelining, every **T/5** seconds, a new instruction completes its **execution**

Design of a Pipeline


* Splitting the Data Path

- * We **divide** the data path into 5 **parts** : IF, OF, EX, MA, and RW

* Timing

- * We insert **latches (registers)** between consecutive stages
- * 4 Latches → IF-OF, OF-EX, EX-MA, and MA-RW
- * At the **negative edge** of a clock, an **instruction** moves from one stage to the next

Pipelined Data Path with Latches




- * Add a latch between subsequent stages.
 - * Triggered by a negative clock edge

The Instruction Packet

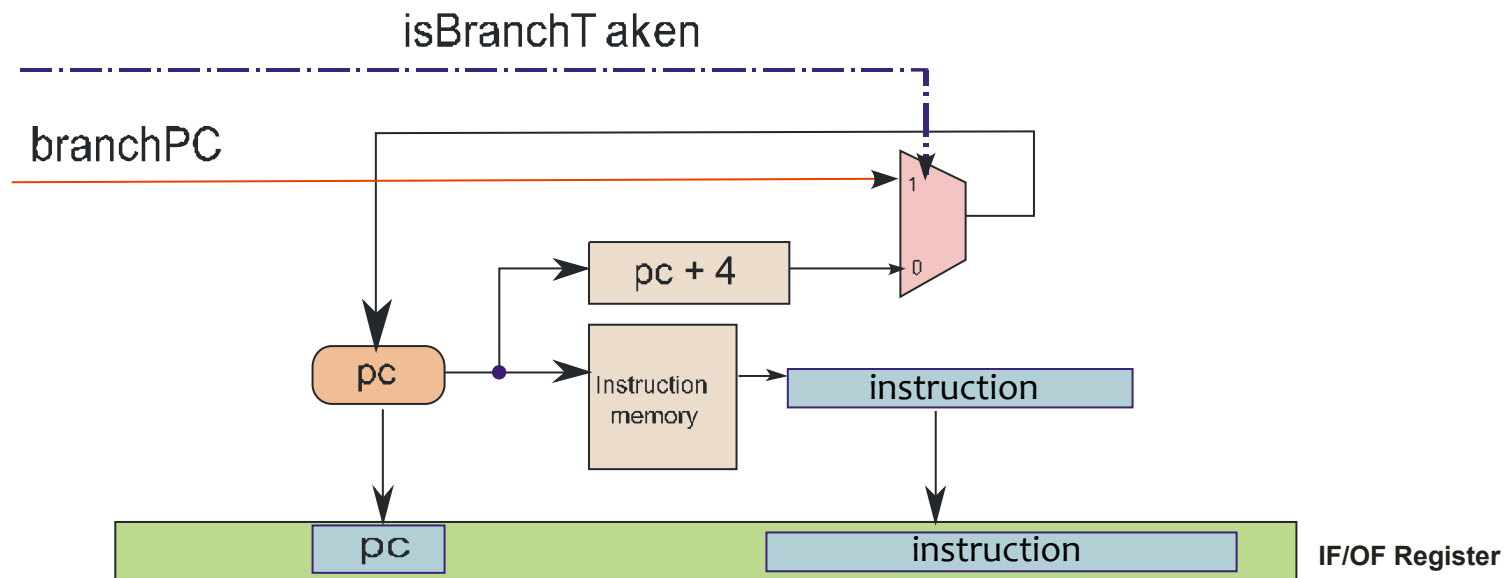
- * What travels **between stages** ?
 - * **ANSWER** : the instruction packet
- * Instruction Packet
 - * Instruction contents
 - * Program counter
 - * All intermediate results
 - * Control signals
- * Every instruction moves with its entire state, no interference between instructions

Outline

- * Overview of Pipelining
- * A Pipelined Data Path
- * Pipeline Hazards
- * Pipeline with Interlocks
- * Forwarding
- * Performance Metrics
- * Interrupts/ Exceptions

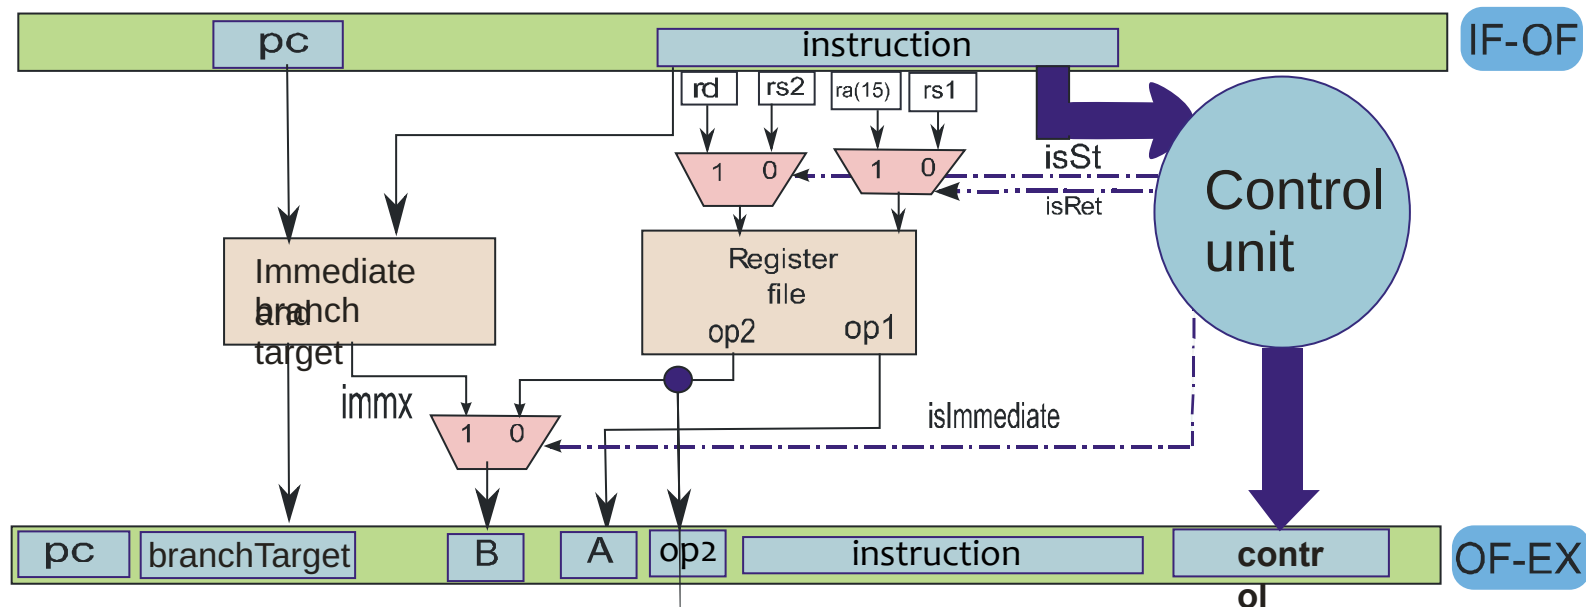


IF Stage



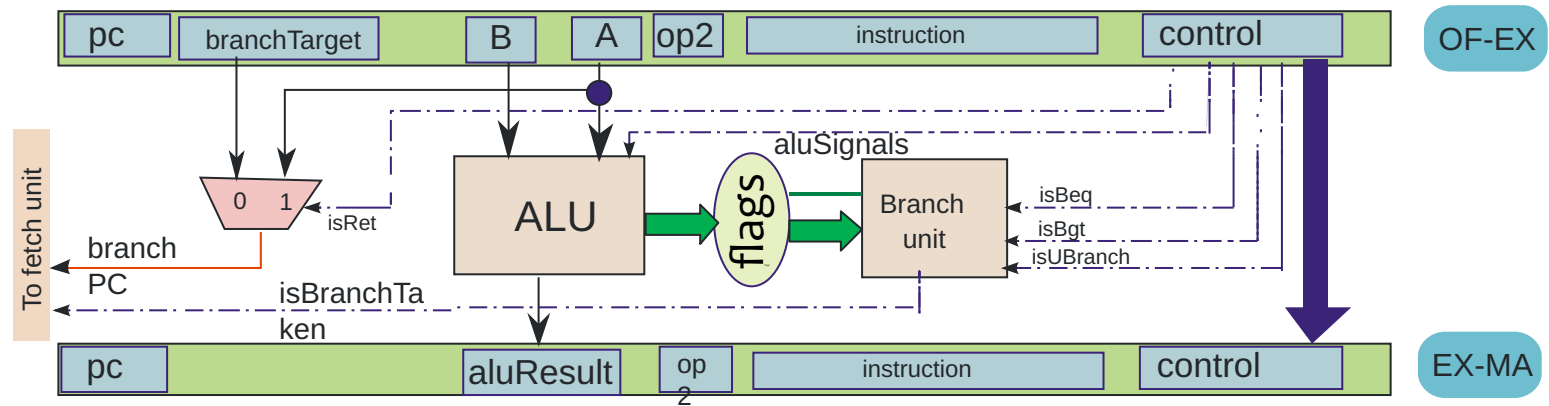
- * Instruction contents saved in the **instruction** field

OF Stage



- * **A, B** → ALU Operands, **op2** (store operand), control (set of all control signals)

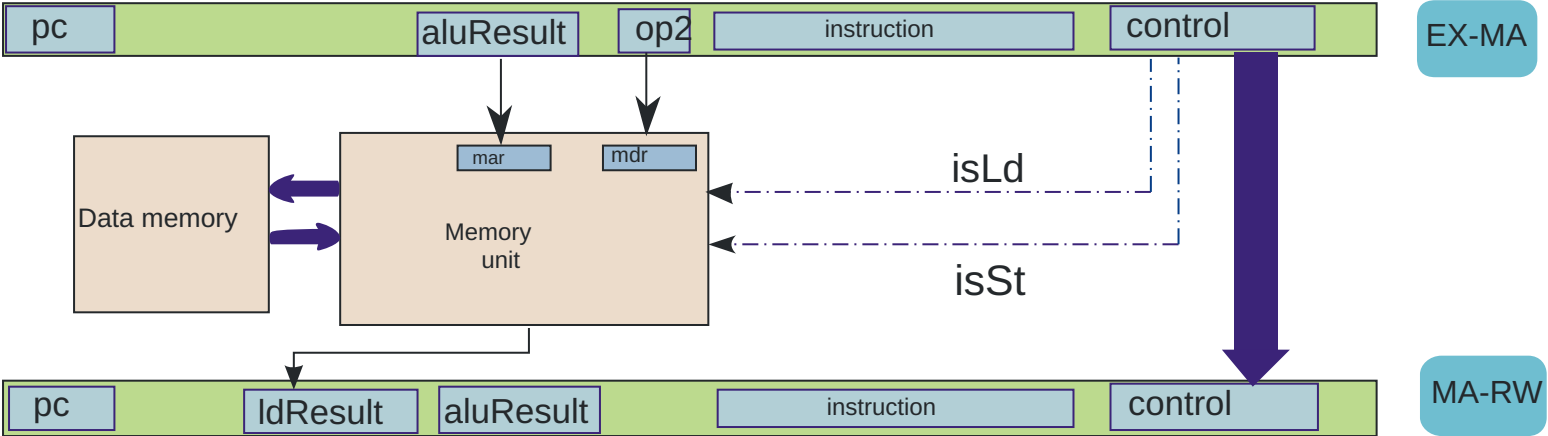
EX Stage



- * **aluResult** → result of the ALU Operation
- * **op2, control, pc, instruction** (passed from OF-EX)

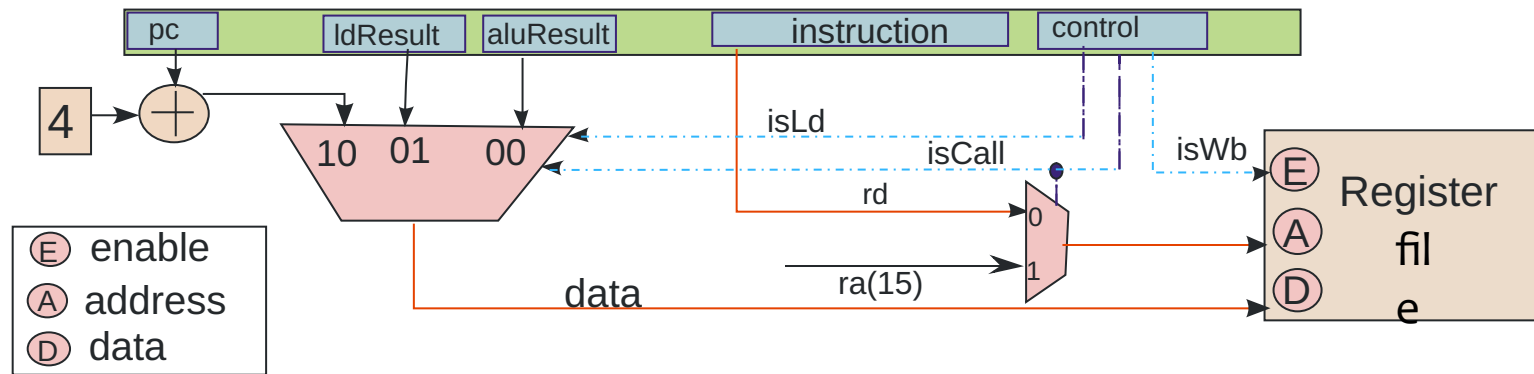
A blue banner with white wavy lines at the bottom. The text "MA Stage" is centered in a large, black, sans-serif font.

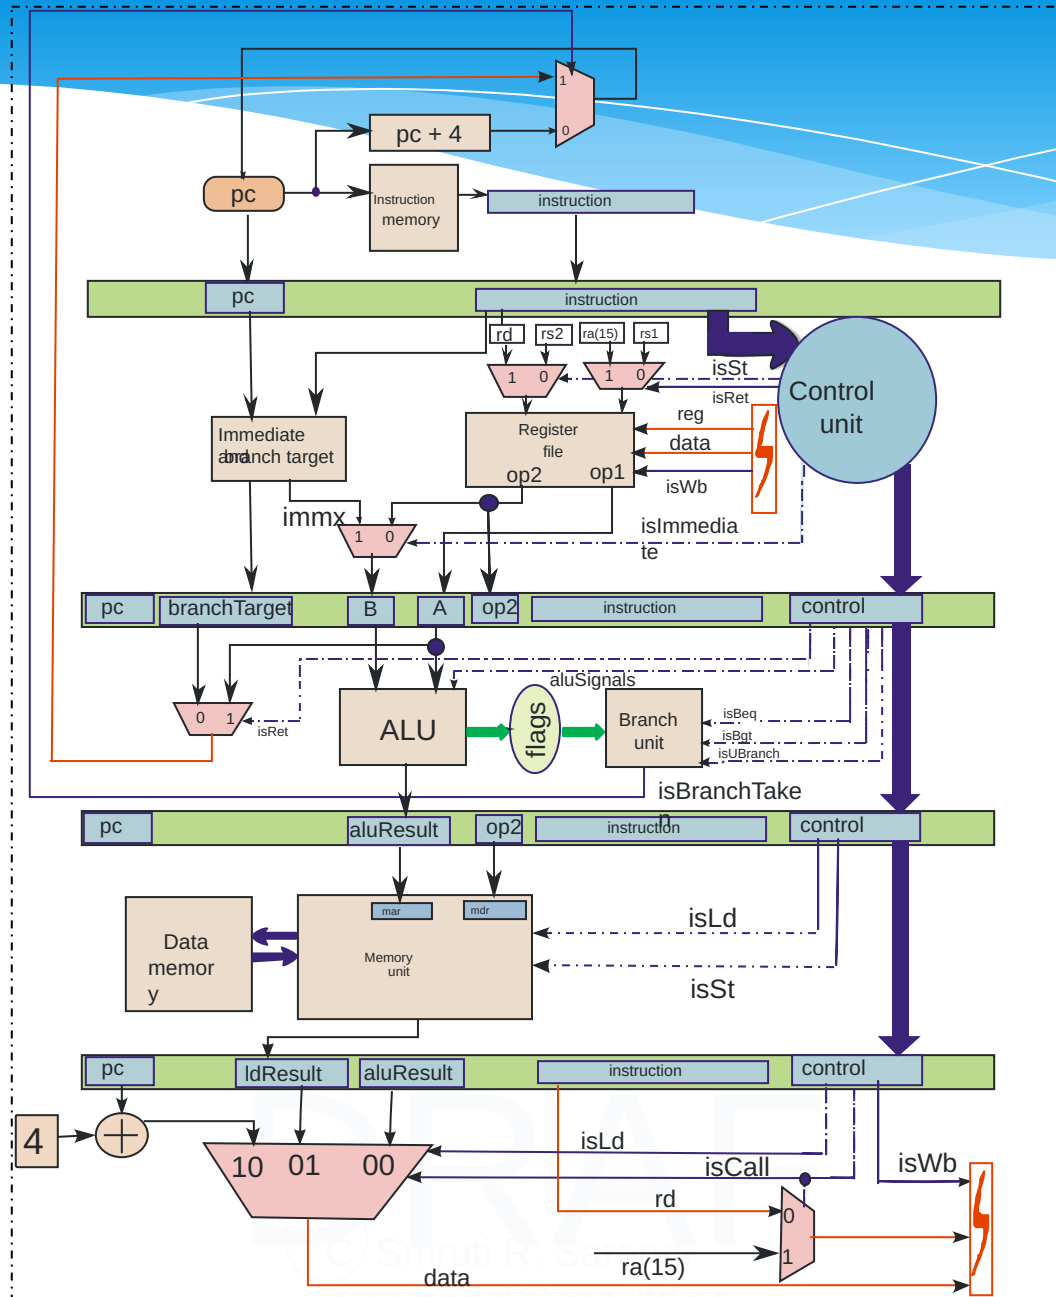
MA Stage



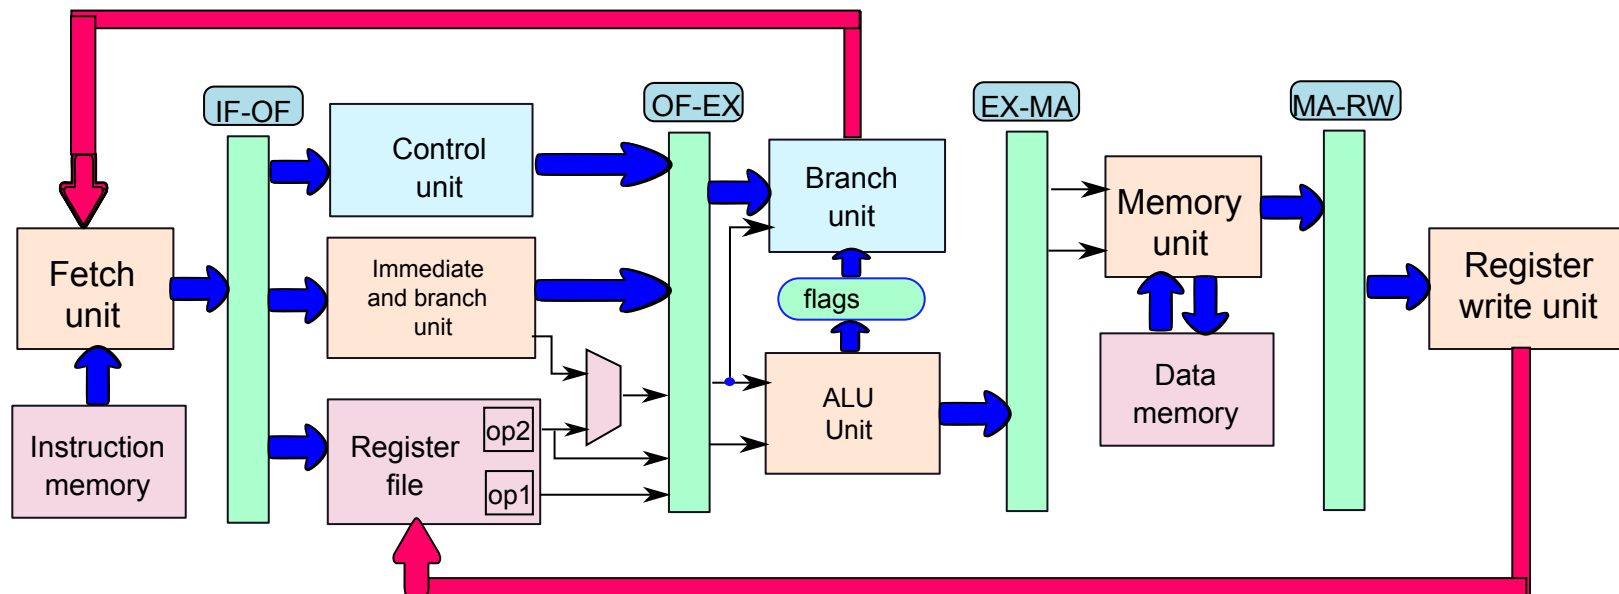
- * **ldResult** → result of the load operation
- * **aluResult, control, pc, instruction** (passed from EX-MA)

RW Stage






Abridged Diagram



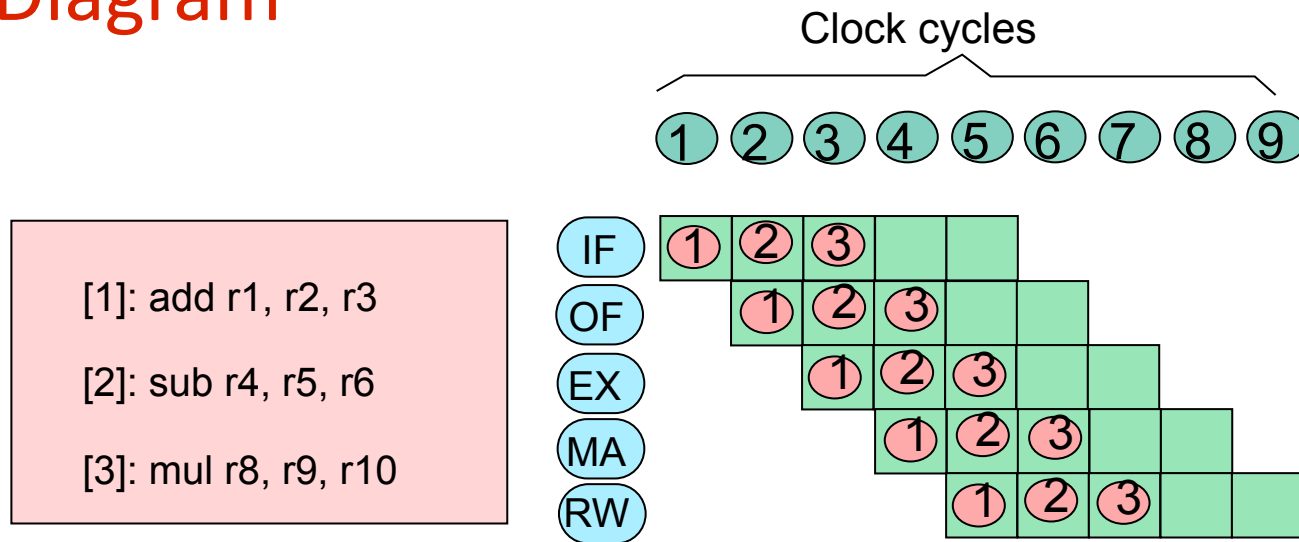
Outline

- * Overview of Pipelining
- * A Pipelined Data Path
- * Pipeline Hazards
- * Pipeline with Interlocks
- * Forwarding
- * Performance Metrics
- * Interrupts/ Exceptions



Pipeline Hazards

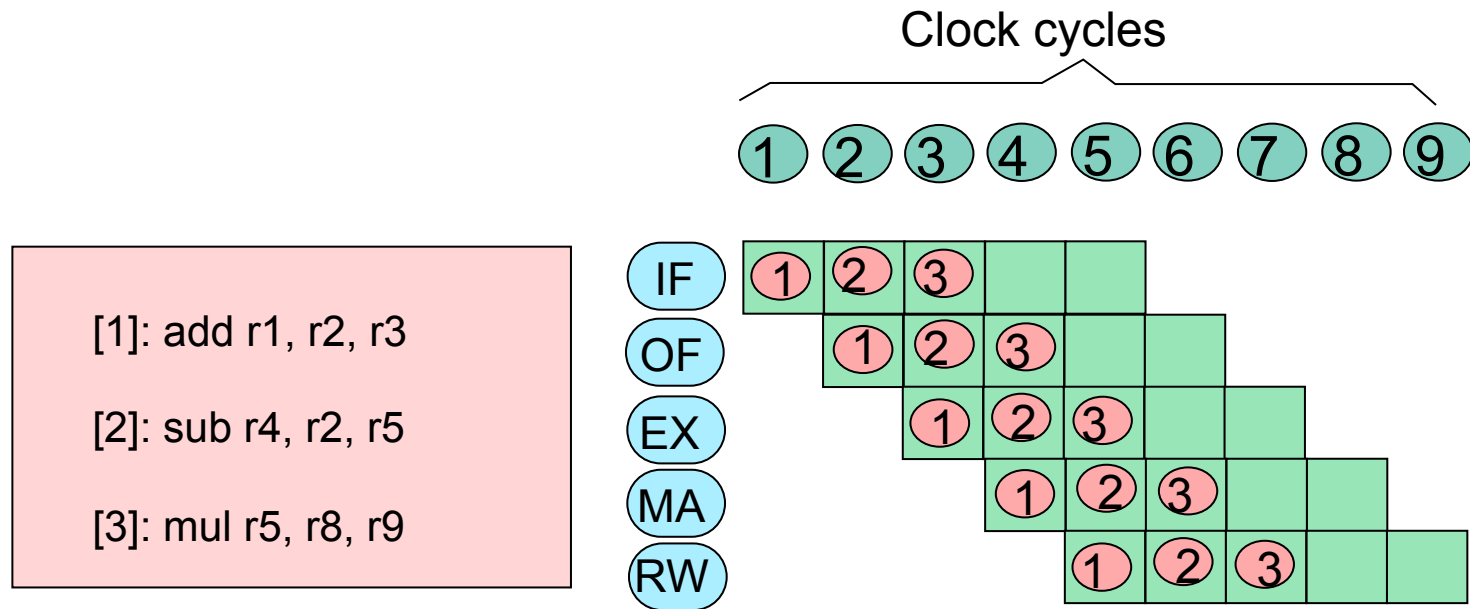
- * Now, let us consider **correctness**
- * Let us introduce a new tool → **Pipeline Diagram**



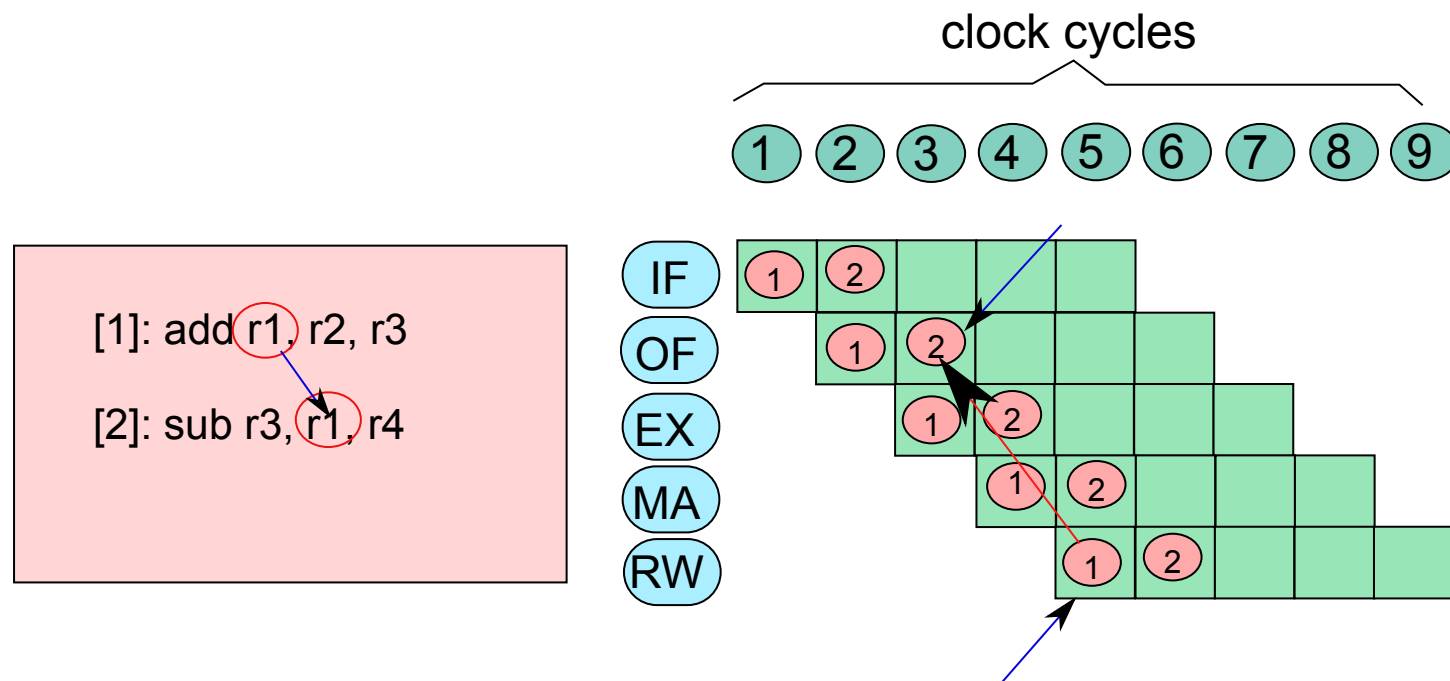
Rules for Constructing a Pipeline Diagram

- * It has 5 **rows**
 - * One per each stage
 - * The rows are **named** : IF, OF, EX, MA, and RW
- * Each **column** represents a clock cycle
- * Each **cell** represents the execution of an instruction in a stage
 - * It is **annotated** with the name(label) of the instruction
- * Instructions proceed from one stage to the next across clock cycles

Example



Data Hazards



* Instruction 2 will read incorrect values !!!

Data Hazard

Definition: A **hazard** is defined as the possibility of erroneous execution of an instruction in a pipeline. A data hazard represents the possibility of erroneous execution because of the unavailability of data, or the availability of **incorrect** data.

- * This situation represents a **data hazard**
- * In specific,
 - * it is a **RAW** (read after write) hazard
- * The earliest we can dispatch instruction 2, is cycle 5